

UE 803 - Data Science

Session 2: Web scraping

Yannick Parmentier - Université de Lorraine / LORIA



Introduction

Web scraping

Automatically extracting data from websites.

Source: [Wikipedia](#)

What is data ?

- audio, image, video, articles, blogs, sensor values, etc.
→ various *data types* (here we focus on **textual data**)
- text files, spreadsheets (e.g. weather forecast), binary files (e.g. word-processing documents), pre-processed data (e.g. treebanks), etc.
→ *(un)structured data* (here we consider **raw data**)
- freely available data (e.g. songs from public domain) vs protected data (e.g. song from a famous artist)
→ *licensed* data (various licenses, to be considered on a case by case basis, here we prioritize **open licenses** such as Creative Commons)

Workflow

Setting up web scrapers can be done in a few steps:

(0. **Target** valuable online data *)

1. **Download** corresponding webpages

2. **Parse** these webpages to extract useful parts

3. **Format and store** the extracted data for further processing

Workflow (continued)

- Note:
 - All this *can be done programmatically* with most major programming languages.
 - Steps 1. and 2. *can be made simpler* if the data host includes a *webservice*
(see next class)

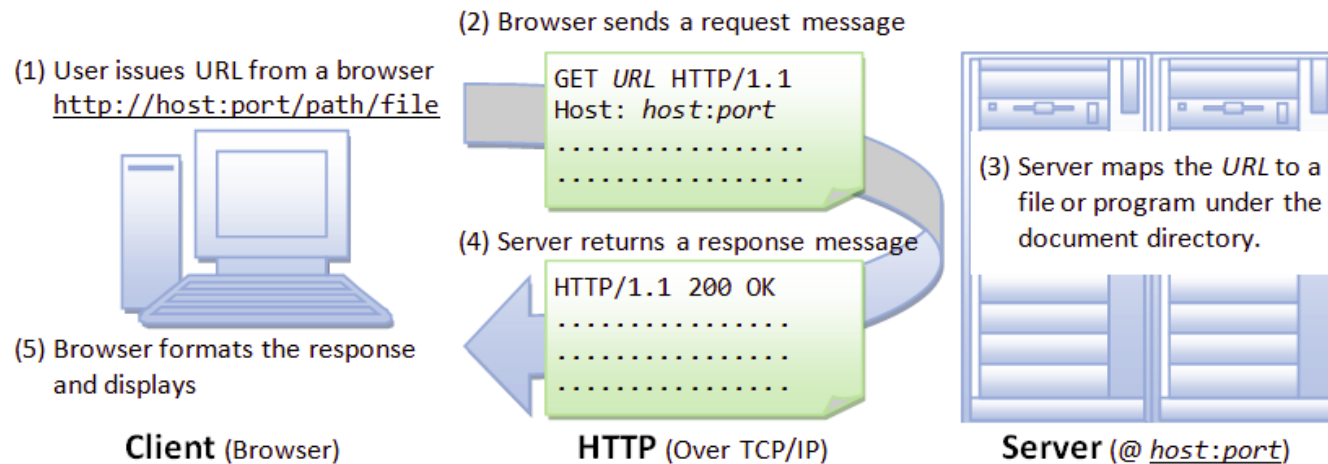
Outline

1. Recall Internet basics (HTTP / HTML)
2. Downloading HTML documents (introducing **requests**)
3. Parsing XML documents (introducing **BeautifulSoup4**)
4. Case study: extracting quotes
5. The automatic way (introducing **Scrapy**)

1. Internet basics

Client/Server architecture

- **Client:** program *requesting* a service (e.g. a *web-browser* such as Firefox requesting access to some file)
- **Server:** program *providing* a service (e.g. a *web-server* such as Apache providing access to some file)
- Communication between these two programs relies on a **communication protocol** (here Hyper Text Transfer Protocol - `http`)



(Picture from F. Boudin's course)

HTTP queries

- Client queries (and server responses) are **lists of bytes** containing a **header** and a **body**
- Main query **types** (*aka* commands):
 - GET → request data from a specified resource (URL)

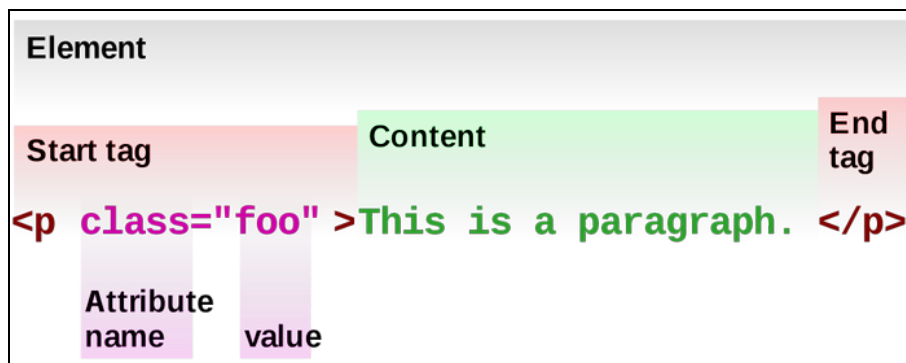
```
GET /test/demo_form.php?name1=value1&name2=value2 HTTP/1.1  
Host: w3schools.com
```

- POST → send data to a server to update a resource

```
POST /test/demo_form.php HTTP/1.1  
Host: w3schools.com  
name1=value1&name2=value2
```

HTML

- Web pages are written in **HyperText Markup Language** (HTML)
- HTML is a **description language**
→ how information *should be* displayed
- HTML is **interpreted** (e.g. by a web-browser)
→ similar to L^AT_EX, python or pdf
- HTML is based on (embedded) **elements** (delimited by opening and closing **tags**, equipped with **attributes**):




(list of HTML [elements](#) and [attributes](#))

HTML (continued)

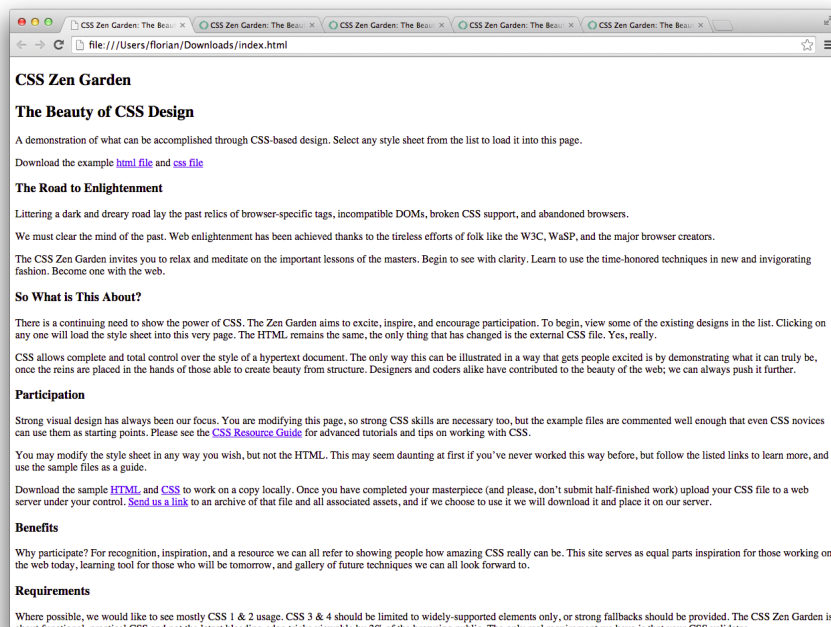
- HTML element *embeddings* follow a **tree structure**:

```
<p>  
  <i> ... </i>  
  <b> ... </b>  
</p>
```



CSS

- Webpages can be associated with **cascading style sheets (CSS)**, which contain **layout definitions**



- Content is thus separated from layout

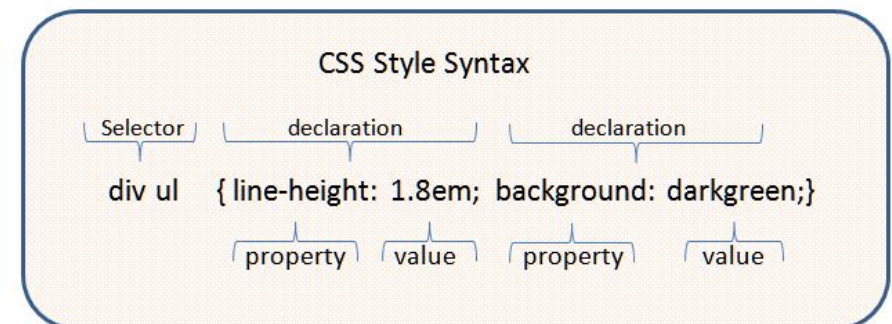
CSS (continued)

- CSS layout definition is declared in the webpage's header:

```
<html lang="en">
<head>
  <meta charset="utf-8">
  <title>Example file</title>

  <link rel="stylesheet" href="/215/215.css">
```

- CSS: list of **rules**
- A rule is made of a **selector** and a list of **declarations** (attribute-values)



CSS selectors

- **Elements** (*, div, ...) or **attributes** (e.g. [width])
- **Dominance** constraints on elements (such as div ul or div > ul)
- **Precedence** constraints on elements (such as div ~ p or div + p)
- **Union** of elements (such as div, ul)

```
<div width="50%">  
...</div>
```

```
<div><ul>  
...  
</ul></div>
```

```
<div>...  
</div>  
<p>...</p>
```

```
<ul>...</ul>
```

```
<div>...</div>
```

CSS selectors (continued)

- **Class** (referring to *groups* of elements, such as `.intro` or `p.intro`)
- **Identifier** (referring to *unique* elements, such as `#intro`)

```
<p class="intro">  
  ...  
  ...  
</p>
```

```
<p id="intro">  
  ...  
  ...  
</p>
```

To inspect a webpage structure : hit the F12 key

The screenshot shows a web browser with multiple tabs. The active tab is 'python4 x' at the URL <https://synalp.loria.fr/python4nlp/>. The website features a large image of a building at night, a navigation bar with links (HOME, PROGRAM, HIGHLIGHTS, REGISTRATION, LOCAL INFORMATION), a 'CONTACT US' link, and a section for 'INVITED SPEAKERS' with two photos. The browser's developer tools are open on the right, showing the 'Elements' panel with the HTML structure of the page. The 'Styles' panel shows the default Bootstrap navbar styles. The 'Network' panel shows a single request to the contact page.

Python4NLP Summer school in Natural Language Processing
26-30 August 2019, Nancy, France

HOME PROGRAM HIGHLIGHTS REGISTRATION LOCAL INFORMATION

CONTACT US

INVITED SPEAKERS

https://synalp.loria.fr/python4nlp/#contact-us

```
...<div class="collapse navbar-collapse" id="bs-example-navbar-collapse-1">
  <ul class="nav navbar-nav navbar-right" id="top-nav">
    <li>
      <a href="#about-us">...</a>
    </li>
    <li>
      <a href="#contact-us">...</a>
    </li>
  </ul>
</div>
```

Styles

```
nav .block .navbar-default
.navbar-nav > li > a {
  color: #2E2E2E;
  padding: 10px 0px;
  margin: 0px 15px;
  font-size: 15px;
  text-transform: uppercase;
}
```


2. Retrieving HTML pages

Introducing requests

Retrieving HTML pages


- Can be done using the `requests` python library:

```
>>> import requests
>>> ua = {'User-agent': 'Mozilla/5.0'}
>>> page=requests.get("http://synalp.loria.fr/index.html", headers=ua)
>>> print(page.status_code)
200
>>> print(page.content)
b'<!DOCTYPE html>\n<html>\n<head>\n  <meta charset="UTF-8">\n
...
```

- Returns an **object** of type `Response` (see [doc.](#) 

```
>>> type(page)
<class 'requests.models.Response'>
```

Retrieving HTML pages (continued)


- NB: The content of the retrieved document (`content` attribute of the `Response` object) is a **binary string**
- If needed, it can be *decoded* (interpreted with a given encoding) using Python's `decode` method (see [doc.](#) 

```
>>> type(page.content)
<class 'bytes'>
>>> print(page.content.decode("utf8"))
<!DOCTYPE html>
<html>
<head>
  <meta charset="UTF-8">
...
```

3. Parsing an XML file

Introducing BeautifulSoup4


Parsing an XML file

- **Parsing** can be done with the BeautifulSoup4 library (see [documentation](#) 

```
>>> from bs4 import BeautifulSoup
>>> soup = BeautifulSoup(page.content, 'html.parser')
>>> print(soup.prettify())
<!DOCTYPE html>
<html>
  <head>
    <meta charset="utf-8"/>
    <meta content="width=device-width, initial-scale=1" name="viewport"/>
  ...
```

- Parsing an XML document amounts to **storing** its content in a **structured** way (i.e. within a dedicated BeautifulSoup object) to make **data access** easier

Parsing an XML file (continued)

- One can **search** through an XML document (that is explore the underlying tree structure, *aka* Document Object Model / DOM) using BeautifulSoup's **dedicated methods** (see [documentation](#) 

```
>>> for para in soup.find_all('p'):
...     print(para.get_text())
...
```

```
Welcome to our website !
The Synalp team is located in Nancy, France and is part of the ...
```

- `find_all(c)` can be used to **extract** the *list* of all **XML elements** (i.e. list of XML **sub-trees**) satisfying a given **constraint** `c`

Searching through an XML document

- Search constraints (**filters**) can be applied:
 - either **on HTML elements**
(via `find`, `find_all`, `find_parent`, etc.)
 - or **on CSS selectors**
(via `select`)
- `find`, `find_next`, etc. retrieve the **first element** satisfying some constraint(s), while
- `find_all`, `find_parents`, etc. retrieve the **list of all elements** satisfying it(them)

Searching through an XML document (continued)

- `find` methods accept constraints on:

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`
 - **attribute name and value**, e.g. `soup.find_all(id='toto')`

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`
 - **attribute name and value**, e.g. `soup.find_all(id='toto')`
 - **element content**, e.g. `soup.find_all(string='toto')`

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`
 - **attribute name and value**, e.g. `soup.find_all(id='toto')`
 - **element content**, e.g. `soup.find_all(string='toto')`
- Note that `find` input parameters can be:

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`
 - **attribute name and value**, e.g. `soup.find_all(id='toto')`
 - **element content**, e.g. `soup.find_all(string='toto')`
- Note that `find` input parameters can be:
 - **list of values**, e.g. `soup.find_all(['a', 'b'])`

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`
 - **attribute name and value**, e.g. `soup.find_all(id='toto')`
 - **element content**, e.g. `soup.find_all(string='toto')`
- Note that `find` input parameters can be:
 - **list of values**, e.g. `soup.find_all(['a','b'])`
 - **regular expressions**, e.g. `soup.find_all(href=re.compile('^https'))`

Searching through an XML document (continued)

- `find` methods accept constraints on:
 - **element name**, e.g. `soup.find_all(name='a')` or `soup.find_all('a')`
 - **attribute name**, e.g. `soup.find_all(href=True)`
 - **attribute name and value**, e.g. `soup.find_all(id='toto')`
 - **element content**, e.g. `soup.find_all(string='toto')`
- Note that `find` input parameters can be:
 - **list of values**, e.g. `soup.find_all(['a','b'])`
 - **regular expressions**, e.g. `soup.find_all(href=re.compile('^https'))`
 - **functions**, e.g.
`soup.find_all(lambda x: x.has_attr('class'))`

Searching through an XML document (continued)

- To return a **limited number of results**: `soup.find_all('a', limit=3)`
and/or to be **non recursive**:
`soup.find_all('a', recursive=False)`

- To access element **contents**:

```
>>> [x.get_text() for x in soup.find_all('a')] #[' Synalp ', ...]
```

- Information can eventually be stored in e.g. **dictionaries** and saved in **files**:

```
>>> data = { 'links' : [x.get_text() for x in soup.find_all('a')] }  
>>> with open('links.txt', 'w') as f:  
...     print(data['links'], file=f)
```

Search examples

```
>>> soup.find('footer') #<footer class="uk-clearfix" ...
>>> soup.find('div', class_='uk-panel') # <div class="uk-panel ...
>>> soup.find('div', attrs={'class': 'uk-panel'})
>>> soup.find_all('a') #[<a class="uk-navbar-brand" ...]
>>> soup.find_all(lambda x: x.has_attr('id'), recursive=False)
>>> soup.select('div p') #[<p class="uk-article-meta"></p>...]
>>> soup.select("p:nth-of-type(3)") #[<p>The Synalp team is ...]
>>> soup.select("div ~ .uk-panel") #[<div class="uk-panel"> <div ...]
```

NB: these retrieve XML elements!

4. Case study

Collect quotes

Finding data

- Say we want to extract quotes from <http://quotes.toscrape.com>:

The screenshot shows a web browser with the URL `quotes.toscrape.com`. The page title is "Quotes to Scrape". The main content area displays three quotes, each with a quote text, author, and tags. The first quote is highlighted with a blue dashed box, and its HTML structure is shown in the browser's developer tools. The developer tools show the following HTML structure for the first quote:

```
<div class="row">
  ::before
  <div class="col-md-8">
    <div class="quote" itemscope="" itemtype="http://schema.org/CreativeWork">
      <span class="text" itemprop="text">
        "The world as we have created it is a process of our thinking. It cannot be changed without changing our thinking."
      </span>
    </div>
  </div>
</div>
```

The developer tools also show the CSS styles for the quote text, including `display: block;`, `margin-bottom: 5px;`, `font-size: large;`, and `font-style: italic;`.

Parsing data

- Useful **HTML element**:

```
<span class="text" itemprop="text">  
    “The world as we have created it is a process of our thinking.  
    It cannot be changed without changing our thinking.”  
</span>
```

- To get *all* the quotes from the page:

```
import requests  
from bs4 import BeautifulSoup  
  
def get_quotes(url):  
    page= requests.get(url=url,headers={'User-Agent': 'Mozilla/5.0'})  
    soup= BeautifulSoup(page.content, 'html.parser')  
    quotes= soup.find_all(class_='text')  
    return list(map(lambda x : x.get_text(), quotes))  
  
url= 'http://quotes.toscrape.com'  
l = get_quotes(url)
```

Parsing data (continued)

- Useful **HTML element** to go to the next page:

```
<li class="next">  
    <a href="/page/2/">Next <span aria-hidden="true">→</span></a>  
</li>
```

- To get the quotes from the next page:

```
def get_next(url):  
    page= requests.get(url=url,headers={'User-Agent': 'Mozilla/5.0'})  
    soup= BeautifulSoup(page.content, 'html.parser')  
    next= soup.find(class_="next")  
    return next  
  
next= get_next(url)  
while next is not None:  
    link= next.find('a')['href']  
    url+= link  
    l.extend(get_quotes(url))  
    next= get_next(url)
```

Saving data

- To store the corresponding data in a text file:

```
with open('quotes.txt','w') as f:  
    for i in range(len(l)):  
        print(i + ": " + l[i], file=f)
```

- Content of the text file (e.g. via `cat quotes.txt`)

```
1: "The world as we have created it is a process of our thinking.  
It cannot be changed without changing our thinking."  
2: "It is our choices, Harry, that show what we truly are, far  
more than our abilities."  
3: "There are only two ways to live your life. One is as though  
nothing is a miracle. The other is as though everything is a  
miracle."  
4: "The person, be it gentleman or lady, who has not pleasure in a  
good novel, must be intolerably stupid."
```

5. The automatic way

Introducing Scrapy

Introducant Scrapy

- open-source framework for web scraping
- includes (among others)
 - a crawling engine,
 - an HTML parser,
 - a JSON exporter
- can be used as a standalone command

Retrieving web pages: spiders

- Defining her.his own spider (class inheriting from Scrapy's Spider class):

```
import scrapy
class QuotesSpider(scrapy.Spider):
```

- A Spider has a unique name, and some start URL(s):

```
class QuotesSpider(scrapy.Spider):
    name = "MyFirstQuoteSpider"
    start_urls = ["http://quotes.toscrape.com/page/1/"]
```

- You can limit scraping to some allowed_domains:

```
allowed_domains=["quotes.toscrape.com"]
```

Parsing configuration

- Apply a `parse` call-back method on each page:

```
class QuotesSpider(scrapy.Spider):  
    name = "MyFirstQuoteSpider"  
    start_urls = ["http://quotes.toscrape.com/page/1/"]  
  
    def parse(self, response):  
        page = response.url.split("/")[-2]  
        filename = "quotes-%s.html" % page  
        with open(filename, "wb") as f:  
            f.write(response.body)
```

- Here, starting pages will be retrieved and stored locally

Running the spider

- Finally, to run this scraping engine, one must use a `CrawlerProcess` as follows:

```
if __name__ == '__main__':  
    import scrapy.crawler  
    myspider = QuotesSpider()  
    process = scrapy.crawler.CrawlerProcess({  
        'USER_AGENT': 'Mozilla/5.0'  
    })  
    process.crawl(myspider)  
    process.start()  
    process.stop()
```

Retrieving and extracting content

- Back to our case study:

The screenshot shows a web browser with multiple tabs. The active tab is 'quotes.toscrape.com/page/1/'. The website has a purple header 'Quotes to Scrape' and a 'Login' link. Below the header, there are three quote cards. The first card features a quote by Albert Einstein: "The world as we have created it is a process of our thinking. It cannot be changed without changing our thinking." with tags: change, deep-thoughts, thinking, world. The second card features a quote by J.K. Rowling: "It is our choices, Harry, that show what we truly are, far more than our abilities." with tags: abilities, choices. The third card features another quote by Albert Einstein: "There are only two ways to live your life. One is as though nothing is a miracle. The other is as though everything is a miracle." with tags: inspirational, life, live, miracle, miracles. A tooltip over the first quote indicates 'span.text | 698 x 50'. The browser's developer tools are open on the right, showing the 'Elements' panel with the HTML structure of the quote cards. The 'Styles' panel shows the CSS for the quote span, including display: block, margin-bottom: 5px, font-size: large, and font-style: italic. The 'Console' panel is empty.

Quotes to Scrape

Login

span.text | 698 x 50

"The world as we have created it is a process of our thinking. It cannot be changed without changing our thinking."

by Albert Einstein (about)

Tags: change deep-thoughts thinking world

"It is our choices, Harry, that show what we truly are, far more than our abilities."

by J.K. Rowling (about)

Tags: abilities choices

"There are only two ways to live your life. One is as though nothing is a miracle. The other is as though everything is a miracle."

by Albert Einstein (about)

Tags: inspirational life live miracle miracles

Elements

```
<div class="row header-box">...</div>
<div class="row">
  ::before
  <div class="col-md-8">
    <div class="quote" itemscope itemtype="http://schema.org/CreativeWork">
      <span class="text" itemprop="text">...</span> == $0
      <span>...</span>
      <div class="tags">...</div>
    </div>
    <div class="quote" itemscope itemtype="http://schema.org/CreativeWork">...</div>
    <div class="quote" itemscope itemtype="http://schema.org/CreativeWork">...</div>
  </div>
</div>
```

html body div.container div.row div.col-md-8 div.quote span.text

Styles

Filter

```
element.style {
}
.quote span.text {
  display: block;
  margin-bottom: 5px;
  font-size: large;
  font-style: italic;
}
* {
  -webkit-box-sizing: border-box;
  -moz-box-sizing: border-box;
  box-sizing: border-box;
}
body {
  font-family: sans-serif;
```

Inherited from body

body { font-family: sans-serif; }

Console

top

Filter

Default levels

Content extraction

- We can extract quotes by using CSS selectors:

```
def parse(self, response):  
    with open('quotes.txt', 'a') as f:  
        quotes = response.css("div.quote")  
        for quote in quotes:  
            text = quote.css('span.text::text').extract_first()  
            print(text, file=f)
```

- Links can be extracted as well:

```
<ul class="pager">  
    <li class="next">  
        <a href="/page/2/">Next (...) </a>  
    </li>  
</ul>
```

Case #1: Extracting links manually

- Using CSS selectors:

```
def parse(self, response):
    with open('quotes.txt', 'a') as f:
        quotes = response.css("div.quote")
        for quote in quotes:
            text = quote.css('span.text::text').extract_first()
            print(text, file=f)
        next_page=response.css('li.next a::attr(href)').extract_first()
        if next_page is not None:
            next_page = response.urljoin(next_page)
            yield scrapy.Request(next_page, callback=self.parse)
```

Case #2: Extracting links using scrapy

- Using the `follow` built-in method:

```
def parse(self, response):
    with open('quotes.txt', 'a') as f:
        quotes = response.css("div.quote")
        for quote in quotes:
            text = quote.css('span.text::text').extract_first()
            print(text, file=f)
        next_page=response.css('li.next a::attr(href)').extract_first()
        if next_page is not None:
            yield response.follow(next_page, callback=self.parse)
```


Thank you!

Slideshow created using **remark**.

Exercise sheet #3 (available on Arche)

- Extracting data from the Internet Movie DataBase